

## Core - Bug # 63615

Bug # 63692 (Accepted): Memory consumption while bulk inserting

<b>Status:</b>	Resolved	<b>Priority:</b>	Must have
<b>Author:</b>	Stephan Großberndt	<b>Category:</b>	Performance
<b>Created:</b>	2014-12-05	<b>Assigned To:</b>	
<b>Updated:</b>	2014-12-10	<b>Due date:</b>	
<b>TYPO3 Version:</b>	6.2		
<b>PHP Version:</b>			
<b>Complexity:</b>	easy		
<b>Is Regression:</b>	No		
<b>Sprint Focus:</b>			
<b>Subject:</b>	Huge memory consumption in DataHandler->processClearCacheQueue while bulk inserting		
<b>Description</b>			
<p>I found a memory consumption problem in      TYPO3\CMS\Core\DataHandling\DataHandler::processClearCacheQueue()</p> <p>due to this dubious construct, which is called for each and every record added:</p>			
<pre>foreach (static::\$recordsToClearCacheFor as \$table =&gt; \$uids) {     foreach (array_unique(\$uids) as \$uid) {         ...          foreach (\$pageldsThatNeedCacheFlush as \$pageld) {             if (\$pageld &gt;= 0) {                 \$tagsToClear[] = 'pageld_' . \$pageld;             }         }         // Queue delete cache for current table and record         \$tagsToClear[] = \$table;         \$tagsToClear[] = \$table . '_' . \$uid;         \$tagsToClear = array_unique(\$tagsToClear);     } }</pre>			
<p>I am adding 6500 records, so <b>FOR EACH AND EVERY</b> record 'pageld_' . \$pageld, \$table and \$table . '_' . \$uid are added to \$tagsToClear and afterwards 'pageld_' . \$pageld and \$table are removed again for 6499 times. There is no reason to do it like this. array_unique() is meant here to reduce memory consumption in \$tagsToClear but for big arrays with array_unique() this is terrible - xhprof showed 2,020,470,496 bytes (~2 gigabyte !!!) of used memory for array_unique() in processClearCacheQueue()</p>			
<p>It is possible to set      clearCache_disable=1</p> <p>to disable clearing of the cache but there is also a simple code solution to solve this since entries in this array are meant to be unique anyway:</p>			
<pre>foreach (static::\$recordsToClearCacheFor as \$table =&gt; \$uids) {     foreach (array_unique(\$uids) as \$uid) {         ...          foreach (\$pageldsThatNeedCacheFlush as \$pageld) {             if (\$pageld &gt;= 0) {                 \$tagsToClear['pageld_' . \$pageld] = 'pageld_' . \$pageld;             }         }     } }</pre>			

```

    }
    // Queue delete cache for current table and record
    $tagsToClear[$table] = $table;
    $tagsToClear[$table . '_' . $uid] = $table . '_' . $uid;
  }
}

```

array\_unique memory consumption drops from

IMemUse%: 13064.1%  
 Incl. MemUse (bytes): 2,020,470,496 bytes (~2 gigabyte !!!)

to

IMemUse%: 4.1%  
 Incl. MemUse (bytes): 629,464 bytes

There may be optimizations possible like

```

foreach (static::$recordsToClearCacheFor as $table => $uids) {
  foreach (array_unique($uids) as $uid) {
    ...

    foreach ($pageIdsThatNeedCacheFlush as $pageId) {
      $pageKey = 'pageId_' . $pageId;
      if ($pageId >= 0 && !isset($tagsToClear[$pageKey])) {
        $tagsToClear[$pageKey] = $pageKey;
      }
    }
    // Queue delete cache for current table and record
    if (!isset($tagsToClear[$table])) {
      $tagsToClear[$table] = $table;
    }
    $tagsToClear[$table . '_' . $uid] = $table . '_' . $uid;
  }
}

```

but even without those its **SOOOO** much better...

Regards,  
 Stephan

## Associated revisions

Revision 0f2f447e - 2014-12-10 01:10 - Stephan Großberndt

[BUGFIX] Memory consumption in DataHandler->processClearCacheQueue

Reduce memory consumption in DataHandler->processClearCacheQueue() for bulk inserts by filling \$tagsToClear as an associative array instead of a normal array with array\_unique()-call after each added record.

Resolves: #63615

Releases: master, 6.2

Change-Id: I1f557f435088cab3aa928ec3f44b0ded9265d7a7

Reviewed-on: <http://review.typo3.org/35085>

Reviewed-by: Lukas Krieger <[lukas.krieger@me.com](mailto:lukas.krieger@me.com)>

Tested-by: Lukas Krieger <[lukas.krieger@me.com](mailto:lukas.krieger@me.com)>

Reviewed-by: Stefan Neufeind <[typo3.neufeind@speedpartner.de](mailto:typo3.neufeind@speedpartner.de)>

Reviewed-by: Alexander Opitz <[opitz.alexander@googlemail.com](mailto:opitz.alexander@googlemail.com)>

Tested-by: Alexander Opitz <[opitz.alexander@googlemail.com](mailto:opitz.alexander@googlemail.com)>

Reviewed-by: Markus Klein <[klein.t3@reelworx.at](mailto:klein.t3@reelworx.at)>

Tested-by: Markus Klein <[klein.t3@reelworx.at](mailto:klein.t3@reelworx.at)>

#### **Revision 31518c96 - 2014-12-10 01:11 - Stephan Großberndt**

[BUGFIX] Memory consumption in DataHandler->processClearCacheQueue

Reduce memory consumption in DataHandler->processClearCacheQueue() for bulk inserts by filling \$tagsToClear as an associative array instead of a normal array with array\_unique()-call after each added record.

Resolves: #63615

Releases: master, 6.2

Change-Id: I1f557f435088cab3aa928ec3f44b0ded9265d7a7

Reviewed-on: <http://review.typo3.org/35211>

Reviewed-by: Markus Klein <[klein.t3@reelworx.at](mailto:klein.t3@reelworx.at)>

Tested-by: Markus Klein <[klein.t3@reelworx.at](mailto:klein.t3@reelworx.at)>

#### **Revision 6da66858 - 2014-12-10 14:34 - Stephan Großberndt**

[BUGFIX] Memory consumption in DataHandler->processClearCacheQueue

Reduce memory consumption in DataHandler->processClearCacheQueue() for bulk inserts by filling \$tagsToClear as an associative array instead of a normal array with array\_unique()-call after each added record.

Resolves: #63615

Releases: master, 6.2

Change-Id: I74bef31fcf1f08665e51cee7efe847e32f501ef2

Reviewed-on: <http://review.typo3.org/35232>

Reviewed-by: Oliver Hader <[oliver.hader@typo3.org](mailto:oliver.hader@typo3.org)>

Tested-by: Oliver Hader <[oliver.hader@typo3.org](mailto:oliver.hader@typo3.org)>

Reviewed-by: Stephan Großberndt <[stephan@grossberndt.de](mailto:stephan@grossberndt.de)>

Tested-by: Stephan Großberndt <[stephan@grossberndt.de](mailto:stephan@grossberndt.de)>

Reviewed-by: Wouter Wolters <[typo3@wouterwolters.nl](mailto:typo3@wouterwolters.nl)>

Tested-by: Wouter Wolters <[typo3@wouterwolters.nl](mailto:typo3@wouterwolters.nl)>

## History

---

### #1 - 2014-12-05 19:28 - Stephan Großberndt

This may be related: <http://forum.typo3.org/index.php?t=msg&#38;th=206887>

### #2 - 2014-12-05 20:59 - Gerrit Code Review

- Status changed from New to Under Review

Patch set 1 for branch **master** of project **Packages/TYPO3.CMS** has been pushed to the review server.

It is available at <http://review.typo3.org/35085>

### #3 - 2014-12-07 11:57 - Gerrit Code Review

Patch set 2 for branch **master** of project **Packages/TYPO3.CMS** has been pushed to the review server.

It is available at <http://review.typo3.org/35085>

### #4 - 2014-12-07 11:59 - Gerrit Code Review

Patch set 3 for branch **master** of project **Packages/TYPO3.CMS** has been pushed to the review server.

It is available at <http://review.typo3.org/35085>

### #5 - 2014-12-07 17:51 - Gerrit Code Review

Patch set 4 for branch **master** of project **Packages/TYPO3.CMS** has been pushed to the review server.

It is available at <http://review.typo3.org/35085>

### #6 - 2014-12-09 07:03 - Philipp Gampe

- Parent task set to #63692

### #7 - 2014-12-09 12:18 - Gerrit Code Review

Patch set 5 for branch **master** of project **Packages/TYPO3.CMS** has been pushed to the review server.

It is available at <http://review.typo3.org/35085>

### #8 - 2014-12-09 12:30 - Gerrit Code Review

*false commit removed*

### #9 - 2014-12-09 12:53 - Gerrit Code Review

*false commit removed*

### #10 - 2014-12-09 15:12 - Lukas Krieger

the gerrit issue-ID is wrong

Gerrit Code Review wrote:

Patch set 3 for branch **master** of project **Packages/TYPO3.CMS** has been pushed to the review server.  
It is available at <http://review.typo3.org/35185>

#### #11 - 2014-12-09 15:28 - Stephan Großberndt

This was an accidental commit to this issue. It should have been #63666 in the first place. It was fixed in the commit message and is shown in the correct issue now.

Stefan Neufeind corrected it here (strike-out and "wrong issue-ID"), but Benjamin Mack did another commit based on the old version. You can safely ignore it here. I requested another delete / strikeout in the typo3.teams.core mailinglist.

#### #12 - 2014-12-09 17:13 - Alexander Opitz

removed false commits here.

Last state is:

Patch set 5 for branch master of project Packages/TYPO3.CMS has been pushed to the review server.  
It is available at <http://review.typo3.org/35085>

#### #13 - 2014-12-10 01:11 - Gerrit Code Review

Patch set 1 for branch **TYPO3\_6-2** of project **Packages/TYPO3.CMS** has been pushed to the review server.  
It is available at <http://review.typo3.org/35211>

#### #14 - 2014-12-10 01:30 - Stephan Großberndt

- Status changed from Under Review to Resolved  
- % Done changed from 0 to 100

Applied in changeset commit:0f2f447e2ec834b69436bd5cca766353bf6d7024.

#### #15 - 2014-12-10 11:26 - Gerrit Code Review

- Status changed from Resolved to Under Review

Patch set 1 for branch **TYPO3\_6-2** of project **Packages/TYPO3.CMS** has been pushed to the review server.  
It is available at <http://review.typo3.org/35232>

#### #16 - 2014-12-10 15:00 - Stephan Großberndt

- Status changed from Under Review to Resolved

